

天氣與人工智慧 II

*Ch.2 The mathematical building blocks of
neural networks*

A first look at a neural network

- A concrete example of a neural network that uses the Python library Keras to learn to classify handwritten digits.
- The problem we're trying to solve here is to classify grayscale images of handwritten digits (28×28 pixels) into their 10 categories (0 through 9).
- We'll use the MNIST dataset, a classic in the machine-learning community. It's a set of 60,000 training images, plus 10,000 test images, assembled by the National Institute of Standards and Technology (the NIST in MNIST) in the 1980s.

A first look at a neural network



Figure 2.1 MNIST sample digits

A first look at a neural network

Listing 2.1 Loading the MNIST dataset in Keras

```
from keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

`train_images` and `train_labels` form the *training set*, the data that the model will learn from. The model will then be tested on the *test set*, `test_images` and `test_labels`.

The images are encoded as Numpy arrays, and the labels are an array of digits, ranging from 0 to 9. The images and labels have a one-to-one correspondence.

A first look at a neural network

Let's look at the training data:

```
>>> train_images.shape
(60000, 28, 28)
>>> len(train_labels)
60000
>>> train_labels
array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

And here's the test data:

```
>>> test_images.shape
(10000, 28, 28)
>>> len(test_labels)
10000
>>> test_labels
array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)
```

A first look at a neural network

- The workflow will be as follows:
 - First, we'll feed the neural network the training data, `train_images` and `train_labels`. The network will then learn to associate images and labels.
 - Finally, we'll ask the network to produce predictions for `test_images`, and we'll verify whether these predictions match the labels from `test_labels`.

A first look at a neural network

- The core building block of neural networks is the *layer*, a data-processing module that you can think of as a filter for data. Some data goes in, and it comes out in a more useful form.
- Most of deep learning consists of chaining together simple layers that will implement a form of progressive *data distillation*.

A first look at a neural network

Listing 2.2 The network architecture

```
from keras import models
from keras import layers
network = models.Sequential()
network.add(layers.Dense(512, activation='relu',
input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))
```

The network consists of a sequence of two Dense layers, which are densely connected (also called *fully connected*) neural layers.

The second (and last) layer is a 10-way *softmax* layer, which means it will return an array of 10 probability scores (summing to 1).

A first look at a neural network

- To make the network ready for training, we need to pick three more things, as part of the *compilation* step:
 - *A loss function*—How the network will be able to measure its performance on the training data, and thus how it will be able to steer itself in the right direction.
 - *An optimizer*—The mechanism through which the network will update itself based on the data it sees and its loss function.
 - *Metrics to monitor during training and testing*—Here, we'll only care about accuracy (the fraction of the images that were correctly classified).

A first look at a neural network

Listing 2.3 The compilation step

```
network.compile(optimizer='rmsprop',  
loss='categorical_crossentropy',  
metrics=['accuracy'])
```

A first look at a neural network

- Before training, we'll preprocess the data by reshaping it into the shape the network expects and scaling it so that all values are in the $[0, 1]$ interval.
- Previously, our training images, for instance, were stored in an array of shape $(60000, 28, 28)$ of type `uint8` with values in the $[0, 255]$ interval. We transform it into a `float32` array of shape $(60000, 28 * 28)$ with values between 0 and 1.
- We also need to categorically encode the labels, a step that's explained in chapter 3.

A first look at a neural network

Listing 2.4 Preparing the image data

```
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype('float32') / 255
```

Listing 2.5 Preparing the labels

```
from keras.utils import to_categorical
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

A first look at a neural network

- We're now ready to train the network, which in Keras is done via a call to the network's fit method—we *fit* the model to its training data:

```
>>> network.fit(train_images, train_labels, epochs=5, batch_size=128)
Epoch 1/5
60000/60000 [=====] - 9s - loss: 0.2524 - acc: 0.9273
Epoch 2/5
51328/60000 [=====>.....] - ETA: 1s - loss: 0.1035 - acc: 0.9692
```

- Two quantities are displayed during training: the loss of the network over the training data, and the accuracy of the network over the training data.

A first look at a neural network

- We quickly reach an accuracy of 0.989 (98.9%) on the training data. Now let's check that the model performs well on the test set, too:

```
>>> test_loss, test_acc =  
network.evaluate(test_images, test_labels)  
>>> print('test_acc:', test_acc)  
test_acc: 0.9785
```

- The test-set accuracy turns out to be 97.8%—that's quite a bit lower than the training set accuracy. This gap between training accuracy and test accuracy is an example of *overfitting*: the fact that machine-learning models tend to perform worse on new data than on their training data. Overfitting is a central topic in chapter 3.

Data representations for neural networks

- In the previous example, we started from data stored in multidimensional Numpy arrays, also called *tensors*.
- In general, all current machine-learning systems use tensors as their basic data structure.
- So what's a tensor?

Scalars (0D tensors)

- A tensor that contains only one number is called a *scalar* (or scalar tensor, or 0-dimensional tensor, or 0D tensor)
- You can display the number of axes of a Numpy tensor via the `ndim` attribute; a scalar tensor has 0 axes (`ndim == 0`). The number of axes of a tensor is also called its *rank*.
- Here's a Numpy scalar:

```
>>> import numpy as np
>>> x = np.array(12)
>>> x
array(12)
>>> x.ndim
0
```


Vectors (1D tensors)

- An array of numbers is called a *vector*, or 1D tensor. A 1D tensor is said to have exactly one axis. Following is a Numpy vector:

```
>>> x = np.array([12, 3, 6, 14])
```

```
>>> x
```

```
array([12, 3, 6, 14])
```

```
>>> x.ndim
```

```
1
```

- This vector has five entries and so is called a *5-dimensional vector*. **Don't confuse a 5D vector with a 5D tensor!** A 5D vector has only one axis and has five dimensions along its axis, whereas a 5D tensor has five axes

Matrices (2D tensors)

- An array of vectors is a *matrix*, or 2D tensor. A matrix has two axes (often referred to *rows* and *columns*). You can visually interpret a matrix as a rectangular grid of numbers. This is a Numpy matrix:

```
>>> x = np.array([[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]])  
>>> x.ndim  
2
```

- The entries from the first axis are called the *rows*, and the entries from the second axis are called the *columns*. In the previous example, [5, 78, 2, 34, 0] is the first row of x, and [5, 6, 7] is the first column.

3D tensors and higher-dimensional tensors

- If you pack such matrices in a new array, you obtain a 3D tensor, which you can visually interpret as a cube of numbers. Following is a

Numpy 3D tensor:

```
>>> x = np.array([[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]],  
[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]],  
[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]]])  
>>> x.ndim  
3
```

key attributes

A tensor is defined by three key attributes:

- *Number of axes (rank)*—For instance, a 3D tensor has three axes, and a matrix has two axes. This is also called the tensor's `ndim` in Python libraries such as Numpy.
- *Shape*—This is a tuple of integers that describes how many dimensions the tensor has along each axis. For instance, the previous matrix example has shape `(3, 5)`, and the 3D tensor example has shape `(3, 3, 5)`. A vector has a shape with a single element, such as `(5,)`, whereas a scalar has an empty shape, `()`.
- *Data type* (usually called `dtype` in Python libraries)—This is the type of the data contained in the tensor; for instance, a tensor's type could be `float32`, `uint8`, `float64`, and so on.

key attributes

To make this more concrete, let's look back at the data we processed in the MNIST example. First, we load the MNIST dataset:

```
from keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

Next, we display the number of axes of the tensor `train_images`, the `ndim`, `shape`, `dtype` attribute:


```
>>> print(train_images.ndim)
3
>>> print(train_images.shape)
(60000, 28, 28)
>>> print(train_images.dtype)
uint8
```

So what we have here is a 3D tensor of 8-bit integers. More precisely, it's an array of 60,000 matrices of 28×28 integers. Each such matrix is a grayscale image, with coefficients between 0 and 255.

Let's display the fourth digit in this 3D tensor, using the library Matplotlib

Listing 2.6 Displaying the fourth digit

```
digit = train_images[4]
import matplotlib.pyplot as plt
plt.imshow(digit,
cmap=plt.cm.binary)
plt.show()
```



Manipulating tensors in Numpy

In the previous example, we *selected* a specific digit alongside the first axis using the syntax `train_images[i]`. Selecting specific elements in a tensor is called *tensor slicing*.

The following example selects digits #10 to #100 (#100 isn't included) and puts them in an array of shape (90, 28, 28):

```
>>> my_slice = train_images[10:100]
>>> print(my_slice.shape)
(90, 28, 28)
```

Manipulating tensors in Numpy

It's equivalent to this more detailed notation, which specifies a start index and stop index for the slice along each tensor axis. Note that `:` is equivalent to selecting the entire axis:

```
>>> my_slice = train_images[10:100, :, :]
>>> my_slice.shape
(90, 28, 28)
>>> my_slice = train_images[10:100, 0:28, 0:28]
>>> my_slice.shape
(90, 28, 28)
```


Manipulating tensors in Numpy

In general, you may select between any two indices along each tensor axis. For instance, in order to select 14×14 pixels in the bottom-right corner of all images, you do this:

```
my_slice = train_images[:, 14:, 14:]
```

It's also possible to use negative indices. Much like negative indices in Python lists, they indicate a position relative to the end of the current axis. In order to crop the images to patches of 14×14 pixels centered in the middle, you do this:

```
my_slice = train_images[:, 7:-7, 7:-7]
```

The notion of data batches

In general, the first axis (axis 0, because indexing starts at 0) in all data tensors you'll come across in deep learning will be the *samples axis* (sometimes called the *samples dimension*). In the MNIST example, samples are images of digits.

In addition, deep-learning models don't process an entire dataset at once; rather, they break the data into small batches. Concretely, here's one batch of our MNIST digits, with batch size of 128:

```
batch = train_images[:128]
```

And here's the next batch:

```
batch = train_images[128:256]
```

Real-world examples of data tensors

Let's make data tensors more concrete with a few examples similar to what you'll encounter later. The data you'll manipulate will almost always fall into one of the following categories:

- *Vector data*—2D tensors of shape (samples, features)
- *Timeseries data or sequence data*—3D tensors of shape (samples, timesteps, features)
- *Images*—4D tensors of shape (samples, height, width, channels) or (samples, channels, height, width)
- *Video*—5D tensors of shape (samples, frames, height, width, channels) or (samples, frames, channels, height, width)

Vector data

This is the most common case. In such a dataset, each single data point can be encoded as a vector, and thus a batch of data will be encoded as a 2D tensor (that is, an array of vectors), where the first axis is the *samples axis* and the second axis is the *features axis*.

Timeseries data or sequence data

Whenever time matters in your data (or the notion of sequence order), it makes sense to store it in a 3D tensor with an explicit time axis. Each sample can be encoded as a sequence of vectors (a 2D tensor), and thus a batch of data will be encoded as a 3D tensor.





Image data

- Images typically have three dimensions: height, width, and color depth.
- Although grayscale images (like our MNIST digits) have only a single color channel and could thus be stored in 2D tensors, by convention image tensors are always 3D, with a one dimensional color channel for grayscale images.
- A batch of 128 grayscale images of size 256×256 could thus be stored in a tensor of shape $(128, 256, 256, 1)$, and a batch of 128 color images could be stored in a tensor of shape $(128, 256, 256, 3)$

Image data



Video data

- Video data is one of the few types of real-world data for which you'll need 5D tensors.
- A video can be understood as a sequence of frames, each frame being a color image.
- Because each frame can be stored in a 3D tensor (`height`, `width`, `color_depth`), a sequence of frames can be stored in a 4D tensor (`frames`, `height`, `width`, `color_depth`), and thus a batch of different videos can be stored in a 5D tensor of shape (`samples`, `frames`, `height`, `width`, `color_depth`).

Video data

- A 60-second, 144×256 YouTube video clip sampled at 4 frames per second would have 240 frames.
- A batch of four such video clips would be stored in a tensor of shape $(4, 240, 144, 256, 3)$.

The gears of neural networks: tensor operations

In our initial example, we were building our network by stacking Dense layers on top of each other. A Keras layer instance looks like this:

```
keras.layers.Dense(512, activation='relu')
```

This layer can be interpreted as a function, which takes as input a 2D tensor and returns another 2D tensor—a new representation for the input tensor. Specifically, the function is as follows (where W is a 2D tensor and b is a vector, both attributes of the layer):

```
output = relu(dot(W, input) + b)
```

Element-wise operations

The `relu` operation and addition are *element-wise* operations: operations that are applied independently to each entry in the tensors being considered. This means these operations are highly amenable to massively parallel implementations.

Element-wise operations

```
def naive_relu(x):  
    assert len(x.shape) == 2  
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] = max(x[i, j], 0)  
    return x
```

x is a 2D Numpy tensor.



Avoid overwriting the input tensor.



Element-wise operations

```
def naive_add(x, y):  
    assert len(x.shape) == 2  
    assert x.shape == y.shape  
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] += y[i, j]  
    return x
```

x and y are 2D Numpy tensors.

Avoid overwriting the input tensor.

Element-wise operations

So, in Numpy, you can do the following element-wise operation:

```
import numpy as np  
z = x + y  
z = np.maximum(z, 0.)
```

Element-wise addition



Element-wise relu

Broadcasting

In the `Dense` layer introduced earlier, we added a 2D tensor with a vector. What happens with addition when the shapes of the two tensors being added differ?

When possible, and if there's no ambiguity, the smaller tensor will be *broadcasted* to match the shape of the larger tensor. Broadcasting consists of two steps:

1. Axes (called *broadcast axes*) are added to the smaller tensor to match the `ndim` of the larger tensor.
2. The smaller tensor is repeated alongside these new axes to match the full shape of the larger tensor.

Broadcasting

Consider X with shape $(32, 10)$ and y with shape $(10,)$.

First, we add an empty first axis to y , whose shape becomes $(1, 10)$.

Then, we repeat y 32 times alongside this new axis, so that we end up with a tensor Y with shape $(32, 10)$,
where $Y[i, :] == y$ for i in $\text{range}(0, 32)$.

At this point, we can proceed to add X and Y , because they have the same shape.

Broadcasting

x is a 2D Numpy tensor.

```
def naive_add_matrix_and_vector(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 1  
    assert x.shape[1] == y.shape[0]  
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] += y[j]  
    return x
```

y is a Numpy vector.

**Avoid overwriting
the input tensor.**

Broadcasting

```
import numpy as np
x = np.random.random((64, 3, 32, 10))
y = np.random.random((32, 10))
z = np.maximum(x, y)
```

```
graph LR
    A["x is a random tensor with shape (64, 3, 32, 10)."] --> B["x = np.random.random((64, 3, 32, 10))"]
    C["y is a random tensor with shape (32, 10)."] --> D["y = np.random.random((32, 10))"]
    B --> E["z = np.maximum(x, y)"]
    D --> E
    E --> F["The output z has shape (64, 3, 32, 10) like x."]
    style A fill:#fde9d9,stroke:#333,stroke-width:1px
    style B fill:none,stroke:none
    style C fill:#fde9d9,stroke:#333,stroke-width:1px
    style D fill:none,stroke:none
    style E fill:none,stroke:none
    style F fill:#fde9d9,stroke:#333,stroke-width:1px
```

x is a random tensor with shape (64, 3, 32, 10).

y is a random tensor with shape (32, 10).

The output z has shape (64, 3, 32, 10) like x.

Tensor dot

- The `dot` operation is the most common, most useful tensor operation.
- An element-wise product is done with the `*` operator in Numpy, Keras, Theano, and TensorFlow.
- `dot` uses a different syntax in TensorFlow, but in both Numpy and Keras it's done using the standard dot operator:

```
import numpy as np
```

```
z = np.dot(x, y)
```

- In mathematical notation, you'd note the operation with a dot (.):

$$z = x \cdot y$$

Tensor dot

```
def naive_vector_dot(x, y):  
    assert len(x.shape) == 1  
    assert len(y.shape) == 1  
    assert x.shape[0] == y.shape[0]  
    z = 0.  
    for i in range(x.shape[0]):  
        z += x[i] * y[i]  
    return z
```

x and y are Numpy vectors.

Tensor dot

```
import numpy as np
def naive_matrix_vector_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0]
    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            z[i] += x[i, j] * y[j]
    return z
```

x is a Numpy matrix.

y is a Numpy vector.

The first dimension of x must be the same as the 0th dimension of y!

This operation returns a vector of 0s with the same shape as y.

Tensor dot

```
def naive_matrix_vector_dot(x, y):  
    z = np.zeros(x.shape[0])  
    for i in range(x.shape[0]):  
        z[i] = naive_vector_dot(x[i, :], y)  
    return z
```

Tensor dot

```
def naive_matrix_dot(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 2  
    assert x.shape[1] == y.shape[0]  
    z = np.zeros((x.shape[0], y.shape[1]))  
    for i in range(x.shape[0]):  
        for j in range(y.shape[1]):  
            row_x = x[i, :]  
            column_y = y[:, j]  
            z[i, j] = naive_vector_dot(row_x, column_y)  
    return z
```

The first dimension of x must be the same as the 0th dimension of y!



The first dimension of x must be the same as the 0th dimension of y!



Tensor dot




Figure 2.5 Matrix dot-product box diagram

Tensor dot

More generally, you can take the dot product between higher-dimensional tensors, following the same rules for shape compatibility as outlined earlier for the 2D case:

$$(a, b, c, d) \cdot (d,) \rightarrow (a, b, c)$$

$$(a, b, c, d) \cdot (d, e) \rightarrow (a, b, c, e)$$

And so on.

Tensor reshaping

We used *tensor reshaping* when we preprocessed the digits data before feeding it into our network:

```
train_images = train_images.reshape((60000, 28 * 28))
```

Reshaping a tensor means rearranging its rows and columns to match a target shape. Naturally, the reshaped tensor has the same total number of coefficients as the initial tensor.

Tensor reshaping

```
>>> x = np.array([[0., 1.],  
                  [2., 3.],  
                  [4., 5.]])  
  
>>> print(x.shape)  
(3, 2)
```

```
>>> x = x.reshape((2, 3))  
>>> x  
array([[ 0.,  1.,  2.],  
       [ 3.,  4.,  5.]])
```

```
>>> x = x.reshape((6, 1))  
>>> x  
array([[ 0.],  
       [ 1.],  
       [ 2.],  
       [ 3.],  
       [ 4.],  
       [ 5.]])
```

```
>>> x = np.zeros((300, 20))  
>>> x = np.transpose(x)  
>>> print(x.shape)  
(20, 300)
```

The engine of neural networks: gradient-based optimization

- Neural networks consist entirely of chains of tensor operations and that all of these tensor operations are just geometric transformations of the input data. It follows that you can interpret a neural network as a very complex geometric transformation in a high-dimensional space, implemented via a long series of simple steps.